

COMPSCI 312 – LAB 2

Marc Thomas

This lab concerns timing C and C++ programs and getting values for both **user** and **system** times. Note that **user** time is the same as **virtual** time and that **profiling** time is the sum of the **user** and **system** times. The **real** time is not very useful because it depends too much on machine load.

1. **First method:** Run your program using the “wrapper:” `/usr/bin/time -p`.

This will run your program using the system calls `fork()` and `execve()` and then print out the **real**, **user**, and **system** times after the program ends. Note that the `-p` option is recommended so that you get the timing information in Posix format. If your program requires user input you can use a Unix pipe. For example, if you want to do a quicksort on the large dataset `datafile6.txt` you could type:

```
sorting datafile6.txt
```

and manually enter the keystrokes “q” for quicksort and “x” to quit. But, you could time this as one atomic operation as follows:

```
echo "q x" | /usr/bin/time -p sorting datafile6.txt
```

The format of the full command line in general will be:

```
echo "{keystrokes}" | /usr/bin/time -p {command with any parameters}
```

Note that:

- i. times reported by the operating system are usually only accurate to $\pm 5\text{--}10\%$, depending on the system, so you may have to do several runs and take an average if you need better accuracy.
- ii. if the program executes with a time under a hundredth of a second you will get 0.00 as an answer, so this method is not useful for quickly executing programs.

Time all of the different sorting options available in the `sorting` program on the large dataset `datafile6.txt`. Are the large variations in time you get surprising?

2. **Second method:** Set timers in your programs, using the timing system calls `setitimer()`, `getitimer()`, etc., or the newer POSIX calls `timer_create()`, `timer_settime()`, etc. To say that these calls are complicated is an *understatement* but I have written the C module `timing.c` with header file `timing.h` which you can use and which is much simpler. To see an example of this, get the sample program `cache.c`. Look at the source code, particularly the lines:

```
#include "timing.h"
...
init_timing(); /* initialize and start timers */
...
read_timing(); /* get the results */
pause_timing(); /* stop the timers */
...
fprintf_timing(); /* write the results to stdout or file */
...
```

Compile the program (with the `timing` module) either via:

```
gcc -g cache.c timing.c -o cache
```

or by using the makefile:

```
make cache
```

`make clean`

This program tests the effective speed of the L2 cache by reading and writing a large amount of information to an array using a particular `line_offset`. It then uses the timing results to compute the net memory bandwidth in megabytes per second for that `line_offset`. Run it, trying some numbers from 127–4096 and see if you can discern a pattern.

Assignment Write a `C` or `C++` program which will accept a positive integer `n` from 2–2000000 on the command line and which will display the first 20 primes greater than or equal to `n`. For an example executable, copy over `cpubound` from the class directory and run it, for example, with `n=2000`:

`cpubound 2000`

Time you program for values of `n=200`, `2000`, `20000` and `200000` using Method 1 (with `/usr/bin/time -p`) above.

Email me your timings (`user` and `system`) for the values of `n` above and the **path** (and **only the path**) to your solution. For example, if your program is named `lab2.c` you might email me that

my solution is at: `/usr/stu/myusername/cs312/lab2.c`

Please:

- i. Keep line length under 80 characters per line.
- ii. Do not send any attachments.